


# Social Media Analytics for Healthcare

# Latent Dirichlet Allocation (LDA)

In natural language processing, **Latent Dirichlet Allocation (LDA)** is a generative statistical model that explains a set of observations through unobserved groups, and each group explains why some parts of the data are similar. LDA is an example of a topic model. In this, observations (e.g., words) are collected into documents, and each word's presence is attributable to one of the document's topics. Each document will contain a small number of topics.



# sklearn.decomposition.LatentDirichletAllocation

```
class sklearn.decomposition.LatentDirichletAllocation(n_components=10, *, doc_topic_prior=None,
topic_word_prior=None, learning_method='batch', learning_decay=0.7, learning_offset=10.0, max_iter=10,
batch_size=128, evaluate_every=-1, total_samples=1000000.0, perp_tol=0.1, mean_change_tol=0.001,
max_doc_update_iter=100, n_jobs=None, verbose=0, random_state=None)
```

[\[source\]](#)

## Parameters:

**n\_components : int, default=10**

Number of topics.

Changed in version 0.19: `n_topics` was renamed to `n_components`

**doc\_topic\_prior : float, default=None**

Prior of document topic distribution `theta`. If the value is None, defaults to `1 / n_components`. In [1], this is called `alpha`.

**topic\_word\_prior : float, default=None**

Prior of topic word distribution `beta`. If the value is None, defaults to `1 / n_components`. In [1], this is called `eta`.

**learning\_method : {'batch', 'online'}, default='batch'**

Method used to update `_component`. Only used in `fit` method. In general, if the data size is large, the online update will be much faster than the batch update.

Valid options:

```
'batch': Batch variational Bayes method. Use all training data in
each EM update.
Old `components_` will be overwritten in each iteration.
'online': Online variational Bayes method. In each EM update, use
mini-batch of training data to update the ``components_``
variable incrementally. The learning rate is controlled by the
``learning_decay`` and the ``learning_offset`` parameters.
```

## Examples


```
>>> from sklearn.decomposition import LatentDirichletAllocation
>>> from sklearn.datasets import make_multilabel_classification
>>> # This produces a feature matrix of token counts, similar to what
>>> # CountVectorizer would produce on text.
>>> X, _ = make_multilabel_classification(random_state=0)
>>> lda = LatentDirichletAllocation(n_components=5,
...     random_state=0)
>>> lda.fit(X)
LatentDirichletAllocation(...)
>>> # get topics for some given samples:
>>> lda.transform(X[-2:])
array([[0.00360392, 0.25499205, 0.0036211 , 0.64236448, 0.09541846],
       [0.15297572, 0.00362644, 0.44412786, 0.39568399, 0.003586  ]])
```

***sklearn.datasets.make\_multilabel\_classification***(n\_samples=100, n\_features=20, n\_classes=5, n\_labels=2, length=50, allow\_unlabeled=True, sparse=False, return\_indicator='dense', return\_distributions=False, random\_state=None)

- ***make\_multilabel\_classification***—Generate a random multilabel classification problem.


# Non-negative matrix factorization

NMF approximates a matrix  $V$  with a low-rank matrix approximation such that  $V \approx W \cdot H$



# Non-negative matrix factorization

- For a given nonnegative matrix  $\mathbf{V}$ , it is possible to find two nonnegative matrices  $\mathbf{W}$  and  $\mathbf{H}$  such that  $\mathbf{V} \approx \mathbf{W} \cdot \mathbf{H}$  holds (ideally).
- In  $\mathbf{V}$ , each column represents an example and each row represents a feature or an attribute.
- $\mathbf{W}$  is called the base matrix and  $\mathbf{H}$  is called the weight matrix or coefficient matrix.



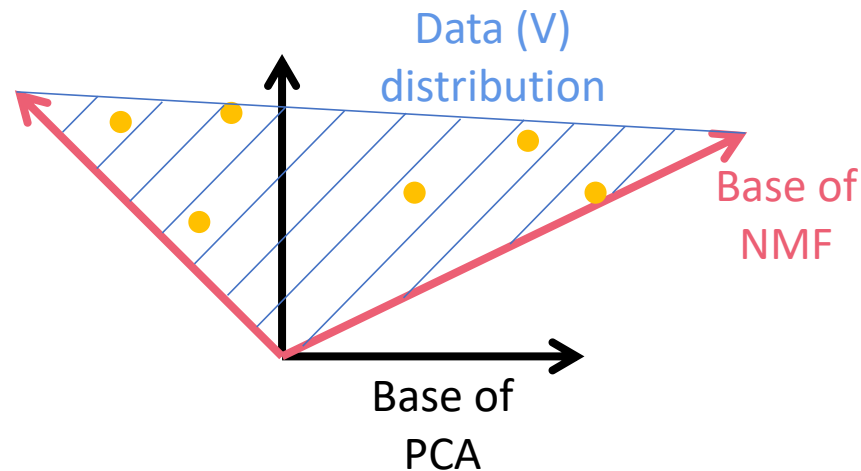
# Non-negative matrix factorization

## Comparison with Principal Component Analysis (PCA):

- PCA:
  - The base vectors are **orthogonal**. They may be distributed inside or outside the data distribution.
  - The reconstruction process simply sums the bases after multiplied by weights, and they do not affect each other.
- NMF:
  - The base vectors are the **edges of** the data **distribution**. They are not necessarily orthogonal.
  - The reconstruction process is realized by computing the matrix inner product.

# Non-negative matrix factorization

Comparison with Principal Component Analysis (PCA):





# Non-negative matrix factorization

## Pros and Cons:

### Pros:

1. Highly interpretable
2. The optimization is convex

Therefore it can be optimized by gradient descent.

### Cons:

1. Factorisation is not unique:

Different optimization paths and optimization stopping times will lead to different results.

2. Feature redundancy among the bases


# Non-negative matrix factorization

## Applications of NMF:

- Compressed Storage
- Information Recovery
- Feature Extraction

## Non-negative signals:

- Image
- Audio
- Natural Language
- Muscle Activity



Using NMF for Image Denoising

# Non-negative matrix factorization

- The main implementation methods of NMF:

## **Bounded optimization**

- The objective of optimization:

$$\operatorname{argmin}_{W,H} \|V - W \cdot H\|_F^2$$

- Major Iterative Methods

- Multiplicative Update Methods

$$W_{ia}^{k+1} = W_{ia}^k \frac{(V(H^k)^T)_{ia}}{(W^k H^k (H^k)^T)_{ia}}, \quad \forall i, a$$

$$H_{bj}^{k+1} = H_{bj}^k \frac{((W^{k+1})^T V)_{bj}}{((W^{k+1})^T W^{k+1} H^k)_{bj}}, \quad \forall b, j.$$

- Gradient Approaches

$$W^{k+1} = \max(0, W^k - \alpha_k \nabla_W f(W^k, H^k)),$$

$$H^{k+1} = \max(0, H^k - \alpha_k \nabla_H f(W^k, H^k)),$$

## sklearn.decomposition.NMF ¶

```
class sklearn.decomposition.NMF(n_components=None, *, init=None, solver='cd', beta_loss='frobenius',
tol=0.0001, max_iter=200, random_state=None, alpha='deprecated', alpha_W=0.0, alpha_H='same',
l1_ratio=0.0, verbose=0, shuffle=False, regularization='deprecated')
```

[\[source\]](#)

Non-Negative Matrix Factorization (NMF).

Find two non-negative matrices, i.e. matrices with all non-negative elements, (W, H) whose product approximates the non-negative matrix X. This factorization can be used for example for dimensionality reduction, source separation or topic extraction.

The objective function is:

$$\begin{aligned} L(W, H) = & 0.5 * ||X - WH||_{loss}^2 \\ & + \alpha_W * l1\_ratio * n\_features * ||vec(W)||_1 \\ & + \alpha_H * l1\_ratio * n\_samples * ||vec(H)||_1 \\ & + 0.5 * \alpha_W * (1 - l1\_ratio) * n\_features * ||W||_{Fro}^2 \\ & + 0.5 * \alpha_H * (1 - l1\_ratio) * n\_samples * ||H||_{Fro}^2 \end{aligned}$$

Where:

$$||A||_{Fro}^2 = \sum_{i,j} A_{ij}^2 \text{ (Frobenius norm)}$$

$$||vec(A)||_1 = \sum_{i,j} abs(A_{ij}) \text{ (Elementwise L1 norm)}$$

The generic norm  $||X - WH||_{loss}$  may represent the Frobenius norm or another supported beta-divergence loss. The choice between options is controlled by the `beta_loss` parameter.

The regularization terms are scaled by `n_features` for W and by `n_samples` for H to keep their impact balanced with respect to one another and to the data fit term as independent as possible of the size `n_samples` of the training set.

The objective function is minimized with an alternating minimization of W and H.

**Parameters:****n\_components : int, default=None**

Number of components, if n\_components is not set all features are kept.

**init : {'random', 'nndsvd', 'nndsvda', 'nndsvdar', 'custom'}, default=None**

Method used to initialize the procedure. Default: None. Valid options:

- `None`: 'nndsvda' if `n_components <= min(n_samples, n_features)`, otherwise random.
- `'random'`: non-negative random matrices, scaled with: `sqrt(X.mean() / n_components)`
- `'nndsvd'`: Nonnegative Double Singular Value Decomposition (NNDSD) initialization (better for sparseness)
- `'nndsvda'`: NNDSD with zeros filled with the average of X (better when sparsity is not desired)
- `'nndsvdar'`: NNDSD with zeros filled with small random values (generally faster, less accurate alternative to NNDSDa for when sparsity is not desired)
- `'custom'`: use custom matrices W and H

*Changed in version 1.1:* When `init=None` and `n_components` is less than `n_samples` and `n_features` defaults to `nndsvda` instead of `nndsvd`.

**solver : {'cd', 'mu'}, default='cd'**

Numerical solver to use: 'cd' is a Coordinate Descent solver. 'mu' is a Multiplicative Update solver.

*New in version 0.17:* Coordinate Descent solver.

*New in version 0.19:* Multiplicative Update solver.

**beta\_loss : float or {'frobenius', 'kullback-leibler', 'itakura-saito'}, default='frobenius'**

Beta divergence to be minimized, measuring the distance between X and the dot product WH. Note that values different from 'frobenius' (or 2) and 'kullback-leibler' (or 1) lead to significantly slower fits. Note that for `beta_loss <= 0` (or 'itakura-saito'), the input matrix X cannot contain zeros. Used only in 'mu' solver.

## Examples

```
>>> import numpy as np
>>> X = np.array([[1, 1], [2, 1], [3, 1.2], [4, 1], [5, 0.8], [6, 1]])
>>> from sklearn.decomposition import NMF
>>> model = NMF(n_components=2, init='random', random_state=0)
>>> W = model.fit_transform(X)
>>> H = model.components_
```